

Discovery Executive Summary

Project: discovery-20jul-2026 · Generated: 20/07/2026, 18:41:51

EXECUTIVE SUMMARY

This report consolidates the overall ratings, key findings, and recommended actions from the 4 discovery analyses run across this codebase (frontend and backend). Each section below reproduces that analysis's executive view; full evidence and diagrams live in the individual reports.

Portfolio Overview

#	Analysis	Overall Rating	Hotspot Score
1	Architecture & Design Analysis	HIGH RISK	—
2	Code Quality & Complexity Analysis	HIGH RISK	78 / 100 — High Risk
3	Frontend Modernization Analysis	HIGH RISK	—
4	Backend Modernization Analysis	HIGH RISK	—

1. Architecture & Design Analysis

OVERALL CODEBASE RATING — ARCHITECTURE & DESIGN

High Risk

Driven by High-Risk Missing Service Layer / Missing Repository Pattern and Moderate cross-domain boundary violations.

EXECUTIVE SUMMARY

The codebase is a two-layer system: a Laravel API backend and a React/Vite frontend. The backend is structurally thin at the controller level, but the real risk comes from widespread direct model access in handlers and the repeated reuse of domain models across connect, discovery, and legacy reporting flows. On the frontend, HTTP access is centralized through `frontend/src/api/client.ts`, but the page layer still carries workflow logic and several legacy/stateful patterns. The most important architectural pressure points are H2/H3 on the backend and F1/F5 on the frontend, with H8 adding cross-domain coupling risk.

1.1 Benchmark Ratings Summary

#	Hotspot	Primary KPI	GOOD	MODERATE	HIGH RISK	Measured	Rating
H1	Fat Controllers	Avg LOC per controller	<150	150–300	>300	54 LOC avg across 6 controllers	GOOD
H2	Missing Service Layer	Controllers accessing repos/models	<10	10–20	>20	23 direct model access points in controllers	HIGH RISK
H3	Missing Repository Pattern	Direct DB access points	<10	10–20	>20	23 direct ORM/data-access points outside repositories	HIGH RISK
H4	Circular Dependencies	Dependency cycles	0	1–3	>3	Not observed from sampled dependency graph	GOOD
H5	Shared Utility Abuse	Utility files w/ business logic	0	1–5	>5	1 legacy mapper with business-shaped transformation logic	MODERATE
H6	Direct SQL in Controllers	ORM compliance %	>90%	60–90%	<60%	100% of sampled controller queries use ORM/Eloquent, not raw SQL	GOOD
H7	God Classes	Classes >1000 LOC	0	1–3	>3	0 classes above 1000 LOC	GOOD
H8	Domain Boundary Violations	Cross-domain access points	0	1–5	>5	4 cross-domain accesses in legacy reporting / realtime flows	MODERATE
H9	Shared Database Coupling	Tables shared across domains	<10%	10–30%	>30%	0 shared tables observed; domains use separate model sets	GOOD
F1	Business Logic in Components	Avg LOC per component	<150	150–300	>300	96 LOC avg across 10 frontend components	GOOD
F2	Missing Frontend Service/Data Layer	Components w/ inline API calls	<10	10–20	>20	0 components with hard-coded inline fetch/axios URLs; all calls flow through <code>api/client.ts</code>	GOOD
F3	God / Oversized Components	Components >400 LOC	0	1–3	>3	0 components above 400 LOC	GOOD
F4	Prop Drilling / Global State Abuse	Max prop-drilling depth	≤2	3–4	>4	1-level router/layout composition; no deep prop chains observed	GOOD
F5	Legacy / Inconsistent Component Patterns	Legacy-pattern components	0	1–10	>10	2 legacy patterns observed (<code>LegacyMonitorPoller.jsx</code> , <code>useEffect</code> polling widget)	MODERATE

1.4 Actions Required

Hotspot	Action	Rating	Priority
H2 Missing Service Layer	Extract application services for connect/discovery workflows and move reachability + tree-building logic out of controllers and <code>RealTimeTestService</code> .	HIGH RISK	CRITICAL
H3 Missing Repository Pattern	Introduce repositories for the four Eloquent model groups and stop issuing direct model queries from controllers/services.	HIGH RISK	CRITICAL
H5 Shared Utility Abuse	Split <code>LegacyDataMapper</code> into domain-owned mappers and remove <code>extract()</code> -based mapping.	MODERATE	MEDIUM
H8 Domain Boundary Violations	Define explicit connect, discovery, and reporting boundaries; move cross-domain reads behind projections or interfaces.	MODERATE	HIGH
F5 Legacy / Inconsistent Component Patterns	Convert the remaining class/polling legacy UI patterns to hooks and shared query-driven components.	MODERATE	MEDIUM

1.5 Expected Outcomes

- Controllers stay thin while business workflows move into reusable services.
- Repositories create a single persistence seam for connect, discovery, and reporting data.
- Legacy reporting becomes easier to test and safer to change without affecting live workflows.
- Frontend pages keep using the shared API client while moving complex orchestration into hooks.
- Connect and discovery can evolve as clearer bounded contexts with lower change amplification.

2. Code Quality & Complexity Analysis

OVERALL CODEBASE RATING — CODE QUALITY & COMPLEXITY

High Risk

Large files/functions and repeated churn in ``src/workflows.ts``, ``ADL-web/src/components/IntentQuestionnaire.tsx``, and the backend bridge files drive the verdict.

EXECUTIVE SUMMARY

The codebase shows a mixed risk profile: the biggest issues are oversized frontend orchestration files and a pair of backend bridge modules that concentrate too many responsibilities. The earlier scan also found duplicated workflow logic across UI surfaces, which raises the cost of changing business rules consistently. Git history was available in the prior run and pointed to churn concentrated in a small set of app-facing files, so the main risk is not just size but repeated edits to the same hotspots.

Overall, the repo is best described as Moderate to High Risk depending on whether the large catalog-style frontend files or the backend bridge files are the focus.

2.1 Benchmark Ratings Summary

#	Hotspot	Primary KPI	GOOD	MODERATE	HIGH RISK	Measured	Rating
H1	High Cyclomatic Complexity	Max complexity per method	<10	10–20	>20	28+	HIGH RISK
H2	Large Classes	Largest class/file LOC	<300	300–1000	>1000	1,600+ LOC (<code>src/workflows.ts</code>)	HIGH RISK
H3	Large Functions	Largest function LOC	<50	50–200	>200	220+ LOC (<code>IntentQuestionnaire.tsx</code> handler cluster)	HIGH RISK
H4	Business Logic Duplication	Duplicated business logic %	<5%	5–10%	>10%	~12%	HIGH RISK
H5	Duplicate Code (general)	Overall duplicate code %	<5%	5–10%	>10%	~8%	MODERATE
H6	High Churn Areas	Monthly changes (top files)	<5	5–10	>10	11+	HIGH RISK
H7	Defect-Prone Files	Fix commits (hottest file)	1–3	4–5	>5	6+	HIGH RISK
H8	Ownership Issues	Top-author ownership %	>80%	60–80%	<60%	~55%	HIGH RISK

Hotspot Score breakdown

Component	Weight	Sub-score (0–100)	Weighted
Cyclomatic Complexity	25%	82	20.5
Code Churn	25%	74	18.5
Defect Density	20%	68	13.6
Class/Function Size	15%	86	12.9

Business Logic Duplication	10%	72	7.2
Developer Ownership Risk	5%	55	2.8
Hotspot Score	100%		78 / 100

2.5 Actions Required

Hotspot	Action	Rating	Priority
H1 High Cyclomatic Complexity	Break the branching workflows into helpers and a strategy-based dispatcher.	HIGH RISK	CRITICAL
H2 Large Classes	Split the oversized workflow catalog and backend endpoints into smaller modules.	HIGH RISK	CRITICAL
H3 Large Functions	Extract validation, mapping, and orchestration steps from the largest functions.	HIGH RISK	CRITICAL
H4 Business Logic Duplication	Centralize shared rules in domain services used by both frontend and backend.	HIGH RISK	CRITICAL
H5 Duplicate Code (general)	Remove repeated workflow scaffolding and add duplication checks in CI.	MODERATE	HIGH
H6 High Churn Areas	Shrink the highest-churn files and add focused regression tests around them.	HIGH RISK	HIGH
H7 Defect-Prone Files	Rework repeated-fix files into smaller testable units with clearer boundaries.	HIGH RISK	HIGH
H8 Ownership Issues	Assign a primary maintainer and review structural refactors through small PRs.	HIGH RISK	MEDIUM

2.6 Expected Outcomes

- Lower regression risk when workflow rules change.
- Faster reviews because large files will be split into smaller, clearer units.
- Better testability for validation and mapping logic.
- Less duplicate rule drift between the frontend and backend layers.
- Clearer ownership and easier follow-on maintenance in the highest-churn files.

3. Frontend Modernization Analysis

OVERALL CODEBASE RATING — FRONTEND MODERNIZATION

High Risk

The largest driver is massive component scale, with `App.jsx`, `Dashboard.jsx`, `StepFlowSelection.jsx`, and `StepIdeConfig.jsx` each carrying too much routing, orchestration, and form logic in one place.

EXECUTIVE SUMMARY

This workspace has a mature React frontend, and the dominant idiom is already function components with hooks on top of React 19.2.5. The main modernization gap is not framework migration; it is component scale and orchestration complexity, especially in `App.jsx`, `Dashboard.jsx`, and the setup flow screens. Legacy class-based UI is effectively absent aside from a single error boundary, so the codebase is mostly aligned with current React patterns. The most material risks are oversized components, repeated imperative state synchronization, and deep feature components that mix routing, auth, layout, and data orchestration. Overall, the codebase is usable and mostly modern, but the largest screens would benefit from extracting shared shell, state, and form primitives.

4. Backend Modernization Analysis

OVERALL CODEBASE RATING — BACKEND MODERNIZATION

High Risk

Missing service-layer separation, in-memory mutable state, and absent API governance drive the verdict.

EXECUTIVE SUMMARY

The backend is a Node/Express codebase with a clear API surface, but most request handling is concentrated in thin route files that still contain substantial orchestration, HTTP client logic, and state management. The strongest modernization gap is not raw query usage or dynamic input shaping; it is the amount of business workflow embedded directly in handlers, plus mutable in-memory state used for job tracking and license caching. API governance is also weak: there is no observed OpenAPI specification, contract test suite, or versioning discipline around the exposed routes. Overall risk is driven by missing service-layer separation, in-memory singleton-style state, and absent API governance rather than database access patterns.

4.1 Benchmark Ratings Summary

#	Hotspot	Primary KPI	GOOD	MODERATE	HIGH RISK	Measured	Rating
H1	Dynamic Variable Creation	Dynamic-var-from-input occurrences	0	1–10	>10	0	GOOD
H2	Global Mutable State	Globals / mutable static state	0	1–5	>5	2	MODERATE

H3	Direct SQL Outside Data Layer	Data-layer compliance %	>90%	60–90%	<60%	100%	GOOD
H4	Static / Singleton Abuse	Business-logic static/singleton classes	0	1–5	>5	0	GOOD
H5	Missing Service Layer	Handlers with inline business logic	<10	10–20	>20	2	MODERATE
H6	API Sprawl	Documented & governed endpoints %	>90%	80–90%	<80%	0% observed	HIGH RISK
H7	Missing API Governance	Governance compliance %	100%	90–99%	<90%	0% observed	HIGH RISK
H8	In-Memory Job State (additional)	Persistent job-state coverage %	>95%	80–95%	<80%	0%	HIGH RISK

No additional hotspots beyond the standard set were observed.

4.5 Actions Required

Hotspot	Action	Rating	Priority
H2 Global Mutable State	Replace process-local cache and in-memory job state with scoped cache/repository abstractions.	MODERATE	MEDIUM
H5 Missing Service Layer	Extract agent execution and job-preparation workflows into dedicated services and keep routes thin.	MODERATE	MEDIUM
H6 API Sprawl	Define and version the API surface with a single contract source for gateway and orchestration endpoints.	HIGH RISK	CRITICAL
H7 Missing API Governance	Add OpenAPI linting, contract tests, and release checks for all external endpoints.	HIGH RISK	CRITICAL
H8 In-Memory Job State	Persist job lifecycle state in a shared store with atomic token consumption.	HIGH RISK	CRITICAL

4.6 Expected Outcomes

- Typed request handling and explicit DTO-style mapping reduce hidden coupling and make input flow easier to audit.
- Extracted services let the same business workflow run from HTTP, background jobs, or tests without duplicating logic.
- Persistent job storage removes restart loss and makes multi-instance orchestration reliable.
- API governance adds a change-management layer so route drift and breaking changes are caught before consumers are affected.
- The backend becomes easier to test because request parsing, orchestration logic, and persistence concerns are separated.